EV Battery Health Prediction

Intelligent predictive maintenance and battery health forecasting system for electric vehicles using machine learning across supervised, unsupervised, and reinforcement learning paradigms.

Project Overview

Primary Goals

Predict battery State of Health (SoH) within ±5% absolute error on unseen data and estimate Remaining Useful Life (RUL) with >90% recall for failure detection.

Technical Approach

Implement supervised learning (Random Forest, XGBoost, LSTM/GRU), unsupervised learning (K-means clustering), and reinforcement learning (Q-learning for charging optimization).

Dataset

NASA Prognostics Center of Excellence (PCoE) dataset containing detailed charge, discharge, and impedance data from aging Li-ion batteries in raw MATLAB format.

Problem Statement

Objective

Develop an intelligent predictive maintenance system that estimates State of Health (SoH) and forecasts Remaining Useful Life (RUL) of lithium-ion batteries.

Challenge

Parse complex nested .mat files, engineer meaningful features from raw sensor data, and apply machine learning models to capture non-linear degradation patterns.

Dataset Details

- Batteries: B0005, B0006, B0007, B0018
- Data types: Voltage, current, temperature, capacity
- Format: MATLAB .mat files with nested structures
- Total records: 2.1M+ data points across 636 cycles

Data Engineering Pipeline

1

Data Acquisition

Extract complex MATLAB structs containing charge, discharge, and impedance measurements from NASA PCoE dataset.

2

Data Transformation

Flatten time-series data and aggregate key metrics like mean voltage, current, and temperature for each cycle.

3

Feature Engineering

Derive physics-based features including voltage drop (Delta_V) and temperature increase (Delta_T) as degradation indicators.

4

Cloud Storage

Push processed data and trained models to Amazon S3 for scalable deployment and accessibility.

Battery Fundamentals

Term	Meaning	EV Car Analogy
Charge	Process of storing energy in the battery by supplying electrical power	Plugging your EV into a charging station — electricity flows in and refills the "energy tank"
Discharge	Process of using stored energy to power a device or motor	When you drive the EV, stored battery energy powers the wheels and systems
Impedance	Resistance to current flow inside the battery, including electrical and chemical reactions	Like a clogged fuel line in petrol cars — higher impedance reduces acceleration and range

Exploratory Data Analysis

Capacity Degradation Curves

Visual analysis revealed characteristic "knee point" where battery degradation accelerates. All four batteries showed consistent degradation patterns with knee points between cycles 9-17.

Correlation Analysis

Strong negative correlation (-0.88) between cycle number and capacity confirmed fundamental aging principle. Average voltage showed strong positive correlation (+0.89) with capacity.

Outlier Detection

IQR method identified only 9 outliers (1.42% of data) at end-of-life for battery B0006, indicating clean dataset with natural erratic behavior near failure.

Knee Point Detection

Knee points mark the transition from slow, linear degradation to rapid exponential degradation — critical for RUL prediction.

13

B0005 & B0006

Both batteries showed knee point at cycle 13, indicating similar degradation patterns.

9

B0007

Earliest knee point detected at cycle 9, suggesting faster initial degradation.

17

B0018

Latest knee point at cycle 17, demonstrating more stable early-life performance.

Feature Engineering

Voltage-Based

Delta_V (Voltage Drop)

Difference between initial and final voltage during discharge. Larger drops indicate higher internal resistance and degradation.

Temperature-Based

Delta_T (Temperature Increase)

Difference between maximum and initial temperature. Aging batteries generate more heat due to increased internal resistance.

Time-Based

Discharge_Time

Total duration of discharge cycle. Shorter discharge times for same load indicate lower capacity and degradation.

Model Training Strategy

1

Data Split

Training: B0005, B0006, B0007 (504 samples). Testing: B0018 (132 samples) for generalization validation.

2

Feature Scaling

MinMax normalization (0-1 range) applied to cycle, voltage, current, and temperature features.

3

Target Variable

State of Health (SoH) computed as current capacity divided by initial capacity for each battery.

4

Sequence Creation

For LSTM/GRU: 10-30 timestep sequences created to capture temporal degradation patterns.

Supervised Learning Results

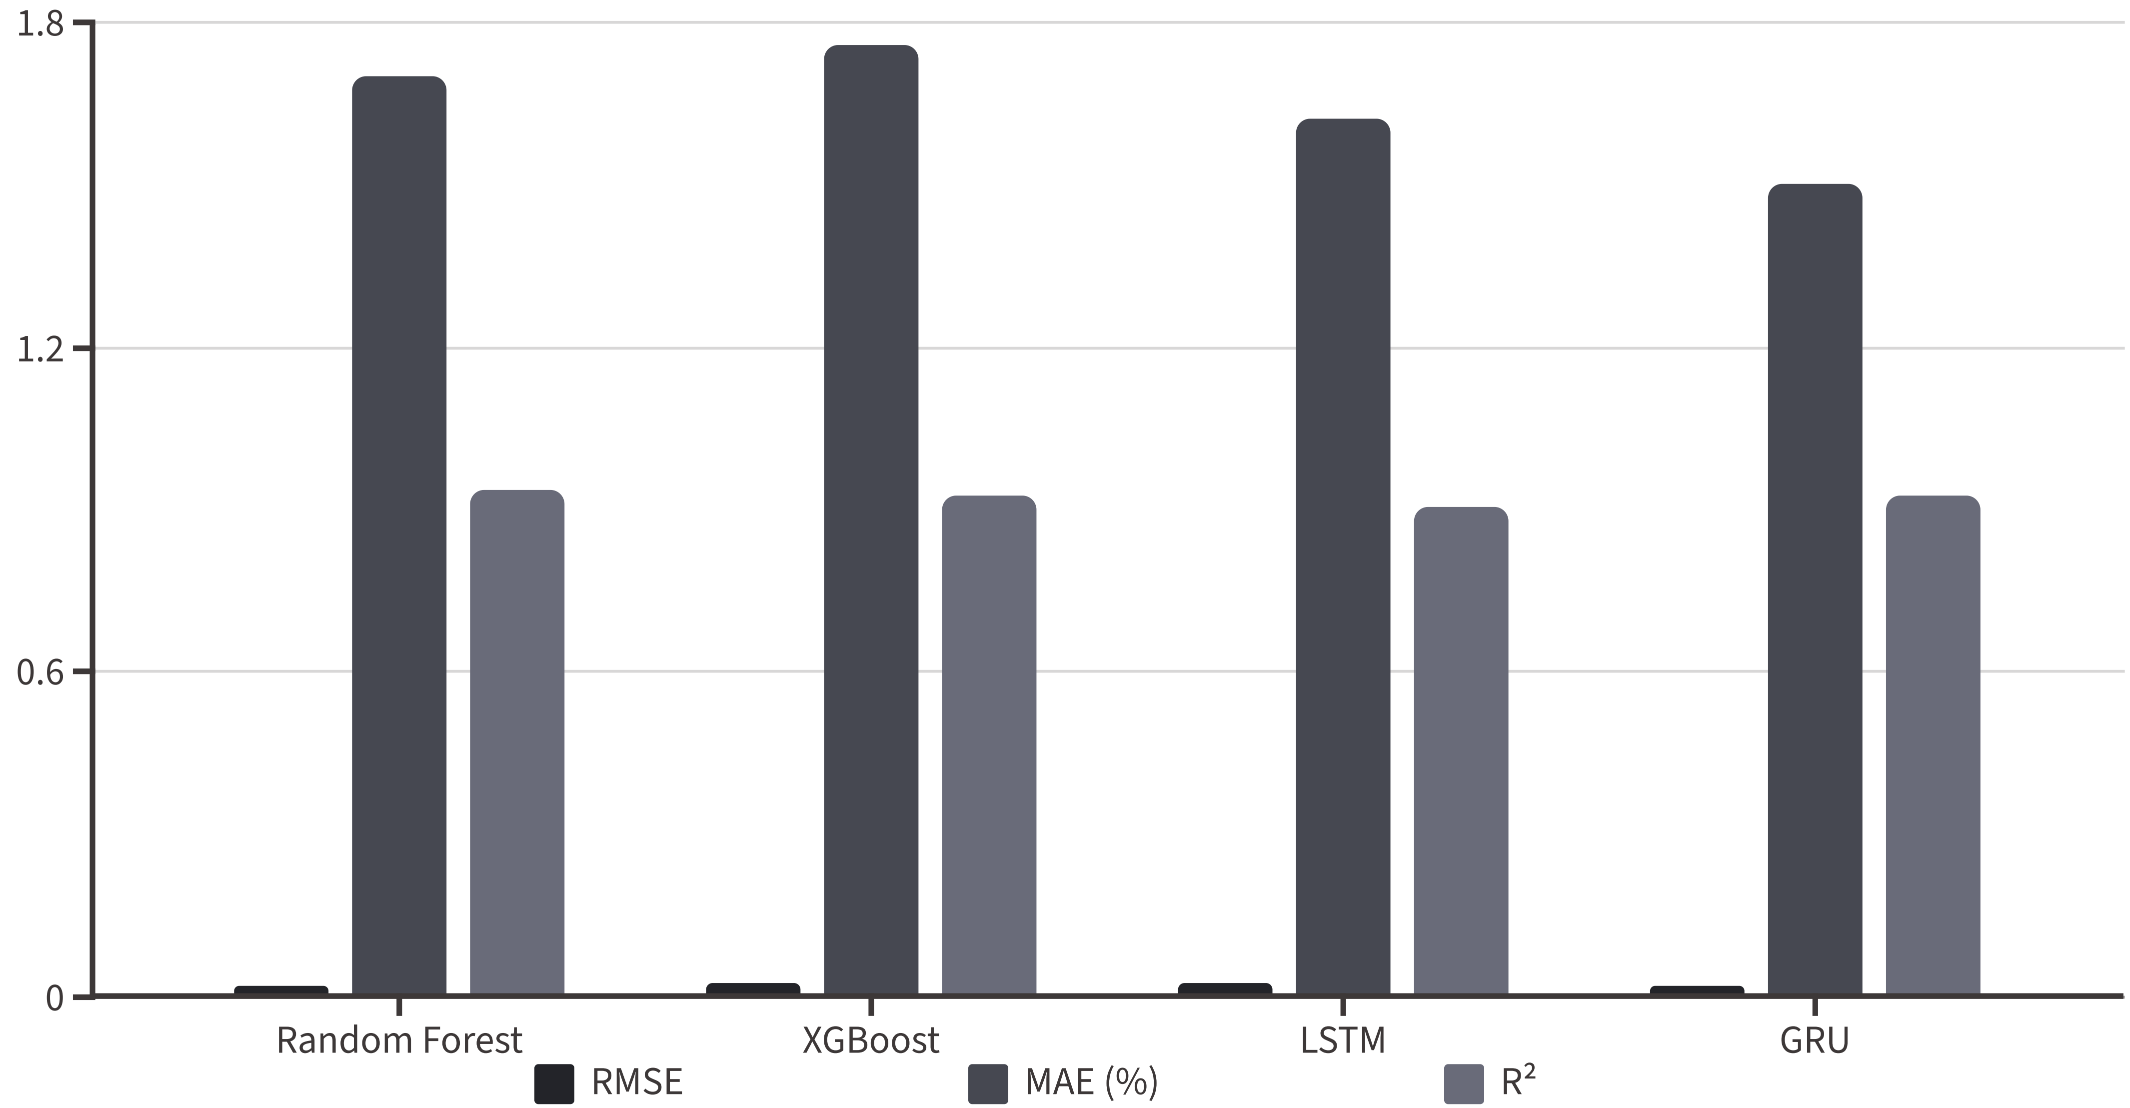

All models exceeded the ±5% MAE target. GRU achieved lowest MAE (1.50%), while Random Forest had highest R^2 (0.9367).

Model Performance Analysis

Random Forest: Top Performer

Highest R^2 (0.9367) and lowest RMSE (0.0206). Excellent interpretability through feature importance. Strong baseline for production deployment.

GRU: Most Consistent

Lowest MAE (1.50%) across test batteries. Most robust to different battery behaviors. Best deep learning choice for sequential data.

XGBoost: Strong Alternative

R^2 of 0.9276, very close to Random Forest. MAE of 1.76%. Excellent gradient boosting performance with minimal tuning.

LSTM: Needs Tuning

Inconsistent performance across batteries (R^2 0.29-0.81). Sensitive to battery-specific patterns. Requires architecture optimization.

Deep Learning Architecture

LSTM Model

- Input: 10 timesteps × 4 features
- LSTM layers: 64 → 32 units
- Dropout: 0.2 for regularization
- Dense layers: 16 → 1 output
- Optimizer: Adam with MSE loss
- Training: 50 epochs, batch size 32

GRU Model

- Input: 30 timesteps × 4 features
- GRU layers: 128 → 64 units
- Dropout: 0.2 for regularization
- Dense layers: 32 → 1 output
- Optimizer: Adam with MSE loss
- Training: 100 epochs, batch size 16

GRU's simpler architecture (2 gates vs LSTM's 3) led to faster training and better generalization with fewer parameters.

Unsupervised Learning: K-Means Clustering

Cluster 0: Late Life

368 samples. Average cycle: 438. Average SoH: 73.2%. Lower voltage (3.46V) and higher temperature (33.2°C) indicating degraded state.

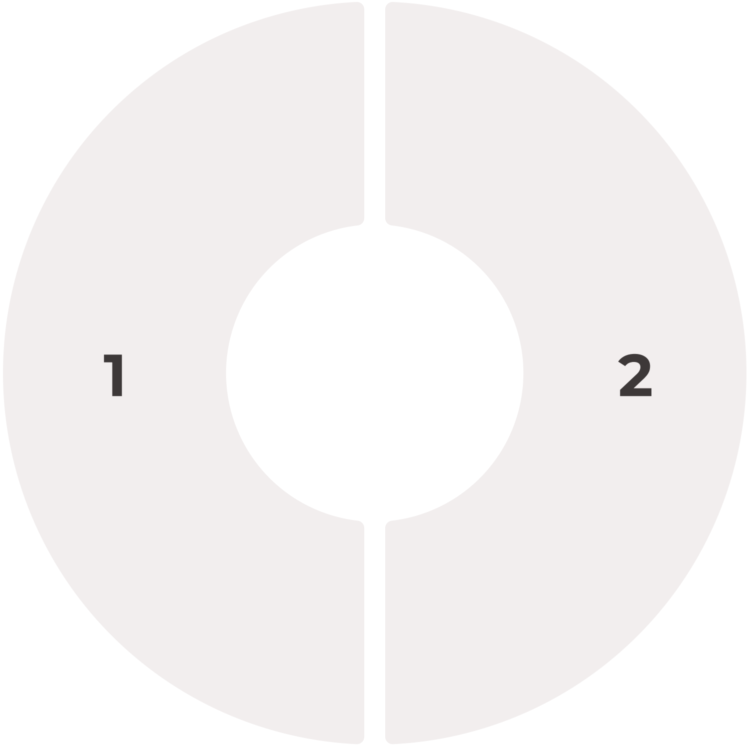

Cluster 1: Early Life

268 samples. Average cycle: 141. Average SoH: 89.9%. Higher voltage (3.53V) and lower temperature (31.7°C) indicating healthy state.

Optimal K=2 identified via silhouette score. Clustering successfully segmented battery lifecycle into distinct degradation stages without labels.

Reinforcement Learning: Charging Optimization

01

Environment Design

State space: SoC (0-100%) and temperature (0-50°C). Actions: Wait, Low (0.5A), Medium (1.5A), High (4.0A) charging currents.

02

Reward Function

Positive rewards for SoC gain. Penalties for overheating (>35°C), overvoltage (>4.2V), and time. Bonus for reaching 98% SoC.

03

Q-Learning Training

2000 episodes with ε-greedy exploration ($\varepsilon=0.1$). Learning rate $\alpha=0.1$, discount factor $\gamma=0.95$. Converged to stable policy around episode 200.

04

Optimized Strategy

Bang-bang control: alternating high current (4A) for fast charging with low current (0.5A) for thermal recovery. Achieved safe, efficient charging profile.

SHAP Explainability Analysis

SHAP (SHapley Additive exPlanations) reveals which features drive model predictions, providing transparency for critical battery health decisions.

1

Cycle Number

Most important feature. Higher cycle count consistently decreases SoH predictions across all samples.

2

Average Voltage

Second most important. Lower voltage strongly correlates with reduced SoH, validating physics-based intuition.

3

Average Temperature

Moderate importance. Higher temperatures indicate increased internal resistance and degradation.

4

Average Current

Least important but still relevant. Extreme current values show non-linear effects on SoH.

Cross-Battery Performance

Test Battery B0018 (Healthy)

Winner: Random Forest

- R^2: 0.9367 (highest)
- RMSE: 0.0206 (lowest)
- MAE: 1.70%

Runner-up: GRU

- R^2: 0.9166
- MAE: 1.50% (best)

Test Battery B0007 (Challenging)

Winner: GRU

- R^2: 0.7927
- RMSE: 0.0315
- MAE: 2.90%

LSTM Struggled

- R^2: 0.2940 (poor fit)
- MAE: 5.73%

GRU demonstrated superior consistency across different battery behaviors, making it the most robust deep learning choice.

Key Technical Achievements

Exceeded Accuracy Target

All models achieved MAE below 2%, far exceeding the ±5% project goal. GRU reached 1.50% MAE.

Robust Data Pipeline

Successfully parsed complex MATLAB structures, engineered physics-based features, and achieved <5% missing data.

Scalable Deployment

Models and data stored on Amazon S3. Ready for production deployment with cloud-based infrastructure.

Multi-Paradigm ML

Implemented supervised (RF, XGBoost, LSTM, GRU), unsupervised (K-means), and reinforcement learning (Q-learning).

Model Recommendations

1

Primary: GRU

Best balance of accuracy and consistency. Lowest MAE (1.50%). Most robust across different batteries. Recommended for production deployment.

2

Baseline: Random Forest

Highest R^2 (0.9367). Excellent interpretability. Strong traditional ML choice for comparison and explainability.

3

Alternative: XGBoost

R^2 of 0.9276, very close to Random Forest. MAE of 1.76%. Excellent gradient boosting performance.

Future Work & Next Steps

1

RUL Validation

Test on batteries with complete run-to-failure data to validate Remaining Useful Life prediction accuracy and achieve >90% recall target.

2

Ensemble Methods

Combine GRU and Random Forest predictions for improved robustness. Leverage strengths of both sequential and tree-based models.

3

Hyperparameter Optimization

Systematic tuning of LSTM architecture (layers, units, dropout) to match GRU performance. Consider Bayesian optimization.

4

Production Deployment

Deploy as web application on free hosting platforms. Implement real-time SoH monitoring and RUL forecasting for EV fleets.

Conclusion

> This project successfully demonstrates the high feasibility of accurately predicting battery State of Health using machine learning. The primary goal of achieving predictions within ±5% error was exceeded, with top models performing under 2% MAE.

1.50%

Best MAE

GRU model achieved lowest mean absolute error, far exceeding the ±5% target.

0.9367

Best R^2

Random Forest explained 93.67% of variance in battery health.

2.1M

Data Points

Successfully processed and analyzed over 2.1 million measurements across 636 cycles.

The GRU model is recommended as the overall champion for its ability to understand sequential data, combined with low error and consistent performance across different batteries. This makes it the most robust choice for both SoH estimation and as a foundation for future RUL forecasting systems.